

debugcm

November 27, 2025

```
[14]: import sys
import numpy as np
import matplotlib.pyplot as plt
import pickle
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torchvision import transforms
from torch.utils.data import Dataset
from torch.utils.data import DataLoader, random_split
import os

np.random.seed(0)
```

```
[15]: if torch.backends.mps.is_available():
    device = torch.device("mps")
    use_mps = True
else:
    device = torch.device("cpu")
    use_mps = False

print(device)
```

mps

```
[16]: class PKLDataset(Dataset):
    def __init__(self, path, transform=None):
        with open(path, "rb") as f:
            data = pickle.load(f)

            self.images = data["images"]          # shape: (N, 28, 28, 3)
            self.labels = data["labels"].reshape(-1)  # shape: (N,) instead of ↵
            ↵ (N, 1)
            self.transform = transform

    def __len__(self):
        return len(self.images)
```

```

def __getitem__(self, idx):
    img = self.images[idx]           # numpy array (28,28,3)
    label = int(self.labels[idx])    # convert to Python int

    # Convert to tensor and permute to (C, H, W)
    img = torch.tensor(img, dtype=torch.float32).permute(2, 0, 1) / 255.0

    if self.transform:
        img = self.transform(img)

    return img, label

```

```
[17]: import pickle

with open("ift-3395-6390-kaggle-2-competition-fall-2025/train_data.pkl", "rb") as f:
    data = pickle.load(f)

print(type(data))
print(len(data) if hasattr(data, "__len__") else "no len")
print(data)
```

```

<class 'dict'>
2
{'images': array([[[[ 6,  4,  0],
   [ 9,  5,  0],
   [ 8,  4,  0],
   ...,
   [ 9,  6,  0],
   [ 9,  6,  0],
   [ 7,  4,  0]],

  [[11,  6,  0],
   [ 4,  4,  0],
   [ 3,  3,  0],
   ...,
   [ 9,  6,  0],
   [ 6,  4,  0],
   [ 4,  2,  0]],

  [[11,  6,  0],
   [ 4,  4,  0],
   [ 3,  3,  0],
   ...,
   [ 6,  4,  0],
   [ 6,  4,  0],
   [ 4,  2,  0]]],
```

```

...,

[[ 1,  1,  0],
 [ 0,  0,  0],
 [ 0,  0,  1],
 ...,
 [ 5,  4,  0],
 [ 6,  5,  0],
 [ 6,  5,  0]],

[[ 3,  1,  1],
 [ 0,  0,  0],
 [ 0,  0,  1],
 ...,
 [ 6,  5,  0],
 [ 6,  5,  0],
 [ 7,  6,  0]],

[[[10,  2,  2],
 [ 0,  0,  1],
 [ 0,  0,  1],
 ...,
 [ 6,  5,  0],
 [ 7,  6,  0],
 [ 7,  6,  0]]],


[[[11,  9,  0],
 [ 9,  7,  0],
 [ 9,  7,  0],
 ...,
 [ 0,  0,  1],
 [ 0,  0,  1],
 [ 0,  0,  1]],

[[12,  9,  0],
 [11,  7,  0],
 [ 9,  6,  0],
 ...,
 [ 0,  0,  2],
 [ 0,  0,  1],
 [ 0,  0,  1]],

[[ 9,  7,  0],
 [12,  8,  0],
 [10,  6,  0],
 ...,

```

```

[ 0,  0,  1],
[ 0,  0,  1],
[ 0,  0,  0]],

...,

[[ 0,  0,  3],
[ 0,  0,  3],
[ 0,  0,  4],
...,

[ 0,  0,  2],
[ 0,  0,  1],
[ 0,  0,  0]],

[[ 0,  0,  2],
[ 0,  0,  2],
[ 0,  0,  5],
...,

[ 0,  0,  2],
[ 0,  0,  1],
[ 1,  0,  0]],

[[ 0,  0,  1],
[ 0,  0,  2],
[ 0,  0,  4],
...,

[ 0,  0,  1],
[ 1,  0,  0],
[ 1,  0,  0]]],


[[[18, 12,  0],
[12,  9,  0],
[17, 11,  0],
...,

[ 0,  0,  3],
[ 5,  0,  2],
[ 5,  0,  2]],

[[15, 11,  0],
[10,  9,  0],
[10,  8,  0],
...,

[ 2,  0,  2],
[ 7,  0,  1],
[ 8,  0,  1]],

[[17, 10,  0],

```

```

[ 7,  6,  0],
[ 4,  4,  0],
...,
[ 0,  0,  2],
[ 5,  0,  1],
[ 8,  0,  1]],

...,

[[ 2,  0,  0],
[ 1,  0,  0],
[ 0,  0,  0],
...,
[ 0,  0,  1],
[ 0,  0,  0],
[ 0,  0,  1]],

[[ 2,  0,  0],
[ 3,  1,  0],
[ 0,  0,  0],
...,
[ 0,  0,  0],
[ 0,  0,  1],
[ 1,  0,  0]],

[[ 3,  0,  0],
[ 2,  0,  0],
[ 2,  1,  0],
...,
[ 0,  0,  1],
[ 1,  0,  0],
[ 1,  0,  0]],

...,

[[[56,  8, 20],
[19,  0, 11],
[ 0,  0,  1],
...,
[ 6,  3,  0],
[11,  7,  0],
[16,  8,  0]],

[[19,  0, 11],
[ 5,  0,  7],
[ 0,  0,  3],
```

```

...,
[ 5,  3,  0],
[ 5,  3,  0],
[ 8,  4,  0]],

[[ 0,  0,  0],
[ 0,  0,  2],
[ 1,  0,  6],
...,
[ 7,  4,  0],
[ 5,  3,  0],
[ 4,  2,  0]],

...,

[[ 4,  3,  0],
[ 4,  2,  0],
[ 4,  2,  0],
...,
[ 0,  0,  1],
[ 1,  0,  0],
[ 0,  0,  0]],

[[ 1,  1,  0],
[ 1,  1,  0],
[ 4,  2,  0],
...,
[ 0,  0,  0],
[ 3,  1,  0],
[ 1,  1,  0]],

[[ 1,  1,  0],
[ 1,  1,  0],
[ 0,  0,  0],
...,
[ 1,  0,  0],
[ 3,  2,  0],
[ 3,  3,  0]]],


[[[ 9,  6,  0],
[ 8,  6,  0],
[ 6,  4,  0],
...,
[ 3,  2,  0],
[ 3,  2,  0],
[ 0,  0,  0]],
```

```

[[ 6,  6,  0],
 [ 4,  4,  0],
 [ 6,  4,  0],
 ...,
 [ 3,  2,  0],
 [ 1,  0,  0],
 [ 2,  0,  0]],

[[ 4,  4,  0],
 [ 6,  4,  0],
 [ 6,  4,  0],
 ...,
 [ 1,  0,  0],
 [ 2,  0,  0],
 [ 3,  0,  0]],

...,

[[ 3,  3,  0],
 [ 3,  3,  0],
 [ 3,  1,  0],
 ...,
 [ 0,  0,  1],
 [ 0,  0,  0],
 [ 0,  0,  0]],

[[ 3,  3,  0],
 [ 3,  3,  0],
 [ 4,  2,  0],
 ...,
 [ 2,  1,  0],
 [ 2,  1,  0],
 [ 1,  1,  0]],

[[ 3,  3,  0],
 [ 3,  3,  0],
 [ 4,  2,  0],
 ...,
 [ 2,  1,  0],
 [ 2,  1,  0],
 [ 1,  1,  0]]],


[[[ 6,  3,  0],
 [ 3,  2,  0],
 [ 2,  0,  2],
 ...,
 [ 7,  6,  0],

```

```

[ 7, 6, 0],
[ 6, 6, 0]],

[[ 4, 2, 0],
[ 1, 0, 1],
[ 2, 0, 2],
...,
[ 8, 5, 0],
[ 7, 5, 0],
[ 6, 4, 0]],

[[ 2, 0, 0],
[ 0, 0, 1],
[ 0, 0, 3],
...,
[ 5, 3, 0],
[ 7, 4, 0],
[ 7, 4, 0]],

...,

[[ 4, 1, 0],
[ 2, 0, 0],
[ 1, 0, 0],
...,
[ 4, 2, 0],
[ 6, 4, 0],
[ 6, 4, 0]],

[[ 7, 3, 0],
[ 7, 4, 0],
[ 6, 2, 0],
...,
[ 6, 4, 0],
[ 7, 4, 0],
[ 7, 4, 0]],

[[11, 6, 0],
[10, 5, 0],
[12, 5, 0],
...,
[ 7, 4, 0],
[ 7, 4, 0],
[ 7, 5, 0]]], shape=(1080, 28, 28, 3), dtype=uint8), 'labels':
array([[0],
       [0],
       [0],
       ...,

```

```

[2],
[2],
[3]], shape=(1080, 1), dtype=uint8) }

[18]: from sklearn.model_selection import train_test_split
from torchvision import transforms
from torch.utils.data import Subset

# Charger le dataset SANS transformation d'abord
dataset = PKLDataset(
    "ift-3395-6390-kaggle-2-competition-fall-2025/train_data.pkl",
)

loader = DataLoader(dataset, batch_size=64, shuffle=False)

d_mean = torch.zeros(3)
d_std = torch.zeros(3)
nb_samples = 0.0

for images, _ in loader:
    batch_samples = images.size(0)

    d_mean += images.mean(dim=[0,2,3]) * batch_samples
    d_std += images.std(dim=[0,2,3]) * batch_samples
    nb_samples += batch_samples

d_mean /= nb_samples
d_std /= nb_samples

d_mean = d_mean.tolist()
d_std = d_std.tolist()

print("Mean:", d_mean)
print("Std:", d_std)

```

Mean: [0.21014535427093506, 0.005330359563231468, 0.2285669893026352]
Std: [0.18871904909610748, 0.01642582379281521, 0.16962255537509918]

```

[19]: # Définir les transformations
transform_train = transforms.Compose([
    transforms.RandomHorizontalFlip(p=0.5), # La rétine n'a pas de sens gauche/
    ↪ droite
    transforms.RandomVerticalFlip(p=0.5), # Ni de haut/bas strict
    transforms.RandomRotation(180),
    transforms.RandomAdjustSharpness(sharpness_factor=2, p=1.0), # La ↪
    ↪ rotation est cruciale pour l'œil

```

```
        transforms.Normalize(mean=d_mean, std=d_std),  
    ])  
  
transform_val = transforms.Compose([  
    transforms.Normalize(mean=d_mean, std=d_std),  
])
```

```
[20]: class TransformSubset(Dataset):  
    def __init__(self, subset, transform=None):  
        self.subset = subset  
        self.transform = transform  
  
    def __getitem__(self, idx):  
        image, label = self.subset[idx]  
        if self.transform:  
            image = self.transform(image)  
        return image, label  
  
    def __len__(self):  
        return len(self.subset)
```

```
[21]: from torch.utils.data import WeightedRandomSampler  
import torch  
import numpy as np  
  
# 1. Votre Split existant (inchangé)  
labels = dataset.labels  
indices = np.arange(len(dataset))  
train_idx, valid_idx = train_test_split(  
    indices,  
    test_size=0.2,  
    random_state=42,  
    stratify=labels  
)  
  
# 2. Préparation du Sampler (NOUVEAU BLOC)  
# On récupère uniquement les labels qui sont dans le set d'entraînement  
y_train = labels[train_idx].reshape(-1) # reshape pour être sûr d'avoir (N,) et non pas (N,1)  
  
# Compter combien d'exemples il y a par classe dans le train  
class_counts = np.bincount(y_train)  
  
# Calculer le poids de chaque classe (Inverse Frequency)  
# Moins la classe est fréquente, plus le poids est grand  
class_weights = 1. / class_counts
```

```

# Assigner un poids à chaque ÉCHANTILLON individuel du train
# Si l'image 1 est de classe 0, elle prend le poids de la classe 0, etc.
samples_weights = class_weights[y_train]
samples_weights = torch.from_numpy(samples_weights).double()

# Créer le sampler
sampler = WeightedRandomSampler(
    weights=samples_weights,
    num_samples=len(samples_weights),
    replacement=True # CRUCIAL : permet de re-piocher les images rares
    ↪plusieurs fois par epoch
)

# 3. Création des Subsets et Transforms (inchangé)
train_data = Subset(dataset, train_idx)
valid_data = Subset(dataset, valid_idx)

train_data = TransformSubset(train_data, transform=transform_train)
valid_data = TransformSubset(valid_data, transform=transform_val)

# 4. DataLoaders (MODIFIÉ)
train_loader = DataLoader(
    train_data,
    batch_size=64,
    sampler=sampler, # <--- On ajoute le sampler ici
    shuffle=False      # <--- OBLIGATOIRE : shuffle doit être False quand on
    ↪utilise un sampler
)

# Le valid_loader reste classique (on ne veut pas de sampler pour la validation)
valid_loader = DataLoader(valid_data, batch_size=128, shuffle=False)

```

```

[22]: import numpy as np

labels = dataset.labels
classes, counts = np.unique(labels, return_counts=True)

from sklearn.utils.class_weight import compute_class_weight

weights = compute_class_weight(
    class_weight="balanced",
    classes=np.unique(labels),
    y=labels
)
class_weights = torch.tensor(weights, dtype=torch.float32).to(device)

loss_fn = nn.CrossEntropyLoss()

```

```

test_loss_fn = nn.CrossEntropyLoss(reduction='sum')

# spot to save your learning curves, and potentially checkpoint your models
savedir = 'results'
if not os.path.exists(savedir):
    os.makedirs(savedir)

```

```

[23]: def train(model, train_loader, optimizer, epoch):
    """Perform one epoch of training."""
    model.train()

    for batch_idx, (inputs, target) in enumerate(train_loader):
        inputs, target = inputs.to(device), target.to(device)

        # 1) Reset gradients
        optimizer.zero_grad()

        # 2) Forward pass
        output = model(inputs)

        # 3) Compute loss
        loss = loss_fn(output, target)

        # 4) Backpropagation
        loss.backward()

        # 5) Update weights
        optimizer.step()

        # Logging
        if batch_idx % 10 == 0:
            print('Train Epoch: {} [{}/{} ({:.0f}%)]\tLoss: {:.6f}'.format(
                epoch,
                batch_idx * len(inputs),
                len(train_loader.dataset),
                100. * batch_idx / len(train_loader),
                loss.item()
            ))

```

```

[24]: from sklearn.metrics import recall_score

def test(model, test_loader):
    """Evaluate the model by doing one pass over a dataset"""
    model.eval()

    test_loss = 0    # total loss over test set

```

```

correct = 0      # total number of correct test predictions
test_size = 0    # number of test samples used
all_preds = []   # to store all predictions
all_targets = [] # to store all targets

with torch.no_grad(): # no backprop, faster evaluation
    for inputs, target in test_loader:
        inputs, target = inputs.to(device), target.to(device)

        # Forward pass
        output = model(inputs)

        # Accumulate loss (sum, not mean)
        loss = test_loss_fn(output, target) # already reduction='sum'
        test_loss += loss.item()

        # Predictions
        pred = output.argmax(dim=1) # index of highest logit

        all_preds.extend(pred.tolist())
        all_targets.extend(target.tolist()) # Target est déjà long pour la suite

        ↪CrossEntropy

        correct += (pred == target).sum().item()

        # Keep track of sample count
        test_size += target.size(0)

# Final metrics
test_loss /= test_size
accuracy = correct / test_size

macro_recall = recall_score(
    all_targets,
    all_preds,
    average='macro',
    zero_division=0
)

print('Test set: Average loss: {:.4f}, Accuracy: {}/{}
      ({:.0f}%)'.format(
    test_loss, correct, test_size, 100. * accuracy))

return test_loss, accuracy, macro_recall, all_preds, all_targets

```

```
[25]: class CNNNet(nn.Module):
    def __init__(self, num_classes=5):
        super().__init__()
```

```

    self.act = nn.ReLU()
    self.drop = nn.Dropout(0.5) # Dropout fort pour éviter le par cœur vu
    ↵qu'on augmente les filtres

    # Bloc 1 : Extraction de features bas niveau (bords, contrastes)
    # On passe de 3 à 64 filtres directement pour capter plus de nuances de
    ↵couleurs
    self.conv1 = nn.Conv2d(3, 64, kernel_size=3, padding=1)
    self.bn1 = nn.BatchNorm2d(64)
    self.conv2 = nn.Conv2d(64, 64, kernel_size=3, padding=1)
    self.bn2 = nn.BatchNorm2d(64)

    # Bloc 2 : Extraction mi-niveau
    self.conv3 = nn.Conv2d(64, 128, kernel_size=3, padding=1)
    self.bn3 = nn.BatchNorm2d(128)
    self.conv4 = nn.Conv2d(128, 128, kernel_size=3, padding=1)
    self.bn4 = nn.BatchNorm2d(128)

    # Bloc 3 : Features complexes
    self.conv5 = nn.Conv2d(128, 256, kernel_size=3, padding=1)
    self.bn5 = nn.BatchNorm2d(256)

    # Classifieur
    # Après 3 max_pool (divisé par 2 trois fois): 28 -> 14 -> 7 -> 3
    self.flatten_dim = 256 * 3 * 3

    self.fc1 = nn.Linear(self.flatten_dim, 512)
    self.fc2 = nn.Linear(512, 5)

def forward(self, x):
    # Bloc 1 (28x28)
    x = self.act(self.bn1(self.conv1(x)))
    x = self.act(self.bn2(self.conv2(x)))
    x = F.max_pool2d(x, 2) # -> 14x14

    # Bloc 2 (14x14)
    x = self.act(self.bn3(self.conv3(x)))
    x = self.act(self.bn4(self.conv4(x)))
    x = F.max_pool2d(x, 2) # -> 7x7

    # Bloc 3 (7x7)
    x = self.act(self.bn5(self.conv5(x)))
    x = F.max_pool2d(x, 2) # -> 3x3

    x = x.view(x.size(0), -1)
    x = self.drop(self.act(self.fc1(x)))
    x = self.fc2(x)

```

```
    return x
```

```
[26]: subset = Subset(train_data, list(range(50)))
subset_loader = DataLoader(subset, batch_size=10, shuffle=False)
model = CNNNet().to(device)
lr = 0.0001
optimizer = optim.Adam(model.parameters(), lr=lr, weight_decay=1e-4)

for epoch in range(50):
    train(model, subset_loader, optimizer, epoch)
    loss, acc, _, _, _ = test(model, subset_loader)
    print(loss, acc)
```

Train Epoch: 0 [0/50 (0%)] Loss: 1.688330

Test set: Average loss: 1.5754, Accuracy: 23/50 (46%)

1.5753730201721192 0.46

Train Epoch: 1 [0/50 (0%)] Loss: 1.387872

Test set: Average loss: 1.5164, Accuracy: 23/50 (46%)

1.5163601684570311 0.46

Train Epoch: 2 [0/50 (0%)] Loss: 1.212225

Test set: Average loss: 1.4519, Accuracy: 25/50 (50%)

1.4518618392944336 0.5

Train Epoch: 3 [0/50 (0%)] Loss: 1.252805

Test set: Average loss: 1.3523, Accuracy: 26/50 (52%)

1.3523392486572265 0.52

Train Epoch: 4 [0/50 (0%)] Loss: 1.185224

Test set: Average loss: 1.2343, Accuracy: 27/50 (54%)

1.23432710647583 0.54

Train Epoch: 5 [0/50 (0%)] Loss: 1.142199

Test set: Average loss: 1.1372, Accuracy: 23/50 (46%)

1.1372257614135741 0.46

Train Epoch: 6 [0/50 (0%)] Loss: 1.492760

Test set: Average loss: 1.0926, Accuracy: 24/50 (48%)

1.092618751525879 0.48

Train Epoch: 7 [0/50 (0%)] Loss: 1.159684

Test set: Average loss: 1.0409, Accuracy: 28/50 (56%)

1.0409052085876465 0.56

Train Epoch: 8 [0/50 (0%)] Loss: 1.249370

Test set: Average loss: 0.9933, Accuracy: 28/50 (56%)

0.9933088207244873 0.56
Train Epoch: 9 [0/50 (0%)] Loss: 1.146795
Test set: Average loss: 0.9777, Accuracy: 30/50 (60%)

0.9776682662963867 0.6
Train Epoch: 10 [0/50 (0%)] Loss: 1.059865
Test set: Average loss: 0.8929, Accuracy: 33/50 (66%)

0.8928742694854737 0.66
Train Epoch: 11 [0/50 (0%)] Loss: 1.199281
Test set: Average loss: 0.9283, Accuracy: 29/50 (58%)

0.9282955741882324 0.58
Train Epoch: 12 [0/50 (0%)] Loss: 1.108542
Test set: Average loss: 0.9264, Accuracy: 31/50 (62%)

0.9263894748687744 0.62
Train Epoch: 13 [0/50 (0%)] Loss: 0.853757
Test set: Average loss: 0.9116, Accuracy: 34/50 (68%)

0.9115592193603516 0.68
Train Epoch: 14 [0/50 (0%)] Loss: 0.907382
Test set: Average loss: 0.8540, Accuracy: 35/50 (70%)

0.8539975833892822 0.7
Train Epoch: 15 [0/50 (0%)] Loss: 0.814083
Test set: Average loss: 0.8630, Accuracy: 35/50 (70%)

0.8630134773254394 0.7
Train Epoch: 16 [0/50 (0%)] Loss: 0.873858
Test set: Average loss: 0.8603, Accuracy: 31/50 (62%)

0.860340461730957 0.62
Train Epoch: 17 [0/50 (0%)] Loss: 0.757625
Test set: Average loss: 0.8772, Accuracy: 34/50 (68%)

0.8771760177612304 0.68
Train Epoch: 18 [0/50 (0%)] Loss: 0.842209
Test set: Average loss: 0.8465, Accuracy: 37/50 (74%)

0.846537799835205 0.74
Train Epoch: 19 [0/50 (0%)] Loss: 0.797163
Test set: Average loss: 0.8708, Accuracy: 31/50 (62%)

0.8708410263061523 0.62
Train Epoch: 20 [0/50 (0%)] Loss: 0.749020
Test set: Average loss: 0.8106, Accuracy: 35/50 (70%)

0.8106286334991455 0.7
Train Epoch: 21 [0/50 (0%)] Loss: 0.767750
Test set: Average loss: 0.8464, Accuracy: 33/50 (66%)

0.8464185810089111 0.66
Train Epoch: 22 [0/50 (0%)] Loss: 0.568536
Test set: Average loss: 0.7560, Accuracy: 35/50 (70%)

0.755951042175293 0.7
Train Epoch: 23 [0/50 (0%)] Loss: 0.609723
Test set: Average loss: 0.7410, Accuracy: 33/50 (66%)

0.7410039710998535 0.66
Train Epoch: 24 [0/50 (0%)] Loss: 0.646144
Test set: Average loss: 0.7125, Accuracy: 36/50 (72%)

0.7124612426757813 0.72
Train Epoch: 25 [0/50 (0%)] Loss: 0.397273
Test set: Average loss: 0.6878, Accuracy: 34/50 (68%)

0.6878188228607178 0.68
Train Epoch: 26 [0/50 (0%)] Loss: 0.426537
Test set: Average loss: 0.6674, Accuracy: 35/50 (70%)

0.6673931503295898 0.7
Train Epoch: 27 [0/50 (0%)] Loss: 0.490698
Test set: Average loss: 0.7037, Accuracy: 39/50 (78%)

0.7037390613555908 0.78
Train Epoch: 28 [0/50 (0%)] Loss: 0.590016
Test set: Average loss: 0.6684, Accuracy: 34/50 (68%)

0.6684346294403076 0.68
Train Epoch: 29 [0/50 (0%)] Loss: 0.407389
Test set: Average loss: 0.7375, Accuracy: 33/50 (66%)

0.7375208377838135 0.66
Train Epoch: 30 [0/50 (0%)] Loss: 0.566266
Test set: Average loss: 0.6601, Accuracy: 38/50 (76%)

0.6600769424438476 0.76
Train Epoch: 31 [0/50 (0%)] Loss: 0.339873
Test set: Average loss: 0.7060, Accuracy: 33/50 (66%)

0.7059865856170654 0.66
Train Epoch: 32 [0/50 (0%)] Loss: 0.483949
Test set: Average loss: 0.7000, Accuracy: 35/50 (70%)

0.6999807929992676 0.7
Train Epoch: 33 [0/50 (0%)] Loss: 0.323711
Test set: Average loss: 0.7188, Accuracy: 35/50 (70%)

0.7187582111358642 0.7
Train Epoch: 34 [0/50 (0%)] Loss: 0.302028
Test set: Average loss: 0.7363, Accuracy: 35/50 (70%)

0.7362630271911621 0.7
Train Epoch: 35 [0/50 (0%)] Loss: 0.269341
Test set: Average loss: 0.6412, Accuracy: 34/50 (68%)

0.6412261390686035 0.68
Train Epoch: 36 [0/50 (0%)] Loss: 0.269679
Test set: Average loss: 0.7475, Accuracy: 33/50 (66%)

0.7475107192993165 0.66
Train Epoch: 37 [0/50 (0%)] Loss: 0.395850
Test set: Average loss: 0.7629, Accuracy: 34/50 (68%)

0.7629396915435791 0.68
Train Epoch: 38 [0/50 (0%)] Loss: 0.366255
Test set: Average loss: 0.6724, Accuracy: 36/50 (72%)

0.6723604106903076 0.72
Train Epoch: 39 [0/50 (0%)] Loss: 0.531890
Test set: Average loss: 0.9044, Accuracy: 34/50 (68%)

0.9044198894500732 0.68
Train Epoch: 40 [0/50 (0%)] Loss: 0.471708
Test set: Average loss: 0.7153, Accuracy: 33/50 (66%)

0.7153428936004639 0.66
Train Epoch: 41 [0/50 (0%)] Loss: 0.296839
Test set: Average loss: 0.6029, Accuracy: 39/50 (78%)

0.6029071426391601 0.78
Train Epoch: 42 [0/50 (0%)] Loss: 0.250599
Test set: Average loss: 0.7204, Accuracy: 33/50 (66%)

0.7204169797897338 0.66
Train Epoch: 43 [0/50 (0%)] Loss: 0.204442
Test set: Average loss: 0.6820, Accuracy: 35/50 (70%)

0.6820449542999267 0.7
Train Epoch: 44 [0/50 (0%)] Loss: 0.157063
Test set: Average loss: 0.7061, Accuracy: 36/50 (72%)

0.7061007785797119 0.72
 Train Epoch: 45 [0/50 (0%)] Loss: 0.280129
 Test set: Average loss: 0.5413, Accuracy: 41/50 (82%)

0.541348533630371 0.82
 Train Epoch: 46 [0/50 (0%)] Loss: 0.159924
 Test set: Average loss: 0.5864, Accuracy: 36/50 (72%)

0.5864129543304444 0.72
 Train Epoch: 47 [0/50 (0%)] Loss: 0.207981
 Test set: Average loss: 0.7096, Accuracy: 34/50 (68%)

0.7095625019073486 0.68
 Train Epoch: 48 [0/50 (0%)] Loss: 0.178077
 Test set: Average loss: 0.7472, Accuracy: 33/50 (66%)

0.747168402671814 0.66
 Train Epoch: 49 [0/50 (0%)] Loss: 0.137306
 Test set: Average loss: 0.6097, Accuracy: 36/50 (72%)

0.6096590662002563 0.72

```
[27]: # TRAINING
model = CNNNet().to(device)
lr=0.001
optimizer = optim.Adam(model.parameters(), lr=lr, weight_decay=1e-4)
scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(
    optimizer, mode='max', factor=0.5, patience=4
)

results = {'name':'model', 'lr': lr, 'train_loss': [],
           'train_acc': [], 'val_loss': [], 'val_acc': [],
           'val_macro_recall': []}
savefile = os.path.join(savedir, results['name']+str(results['lr'])+'.pkl' )

for epoch in range(1, 50):
    train(model, train_loader, optimizer, epoch)

    train_loss, train_acc, _, _, _ = test(model, train_loader)
    val_loss, val_acc, val_macro_recall, final_val_preds, final_val_targets = test(model, valid_loader)

    scheduler.step(val_acc)
```

```

results['train_loss'].append(train_loss)
results['train_acc'].append(train_acc)

results['val_loss'].append(val_loss)
results['val_acc'].append(val_acc)

results['val_macro_recall'].append(val_macro_recall)

if epoch == 49:
    results['final_val_preds'] = final_val_preds
    results['final_val_targets'] = final_val_targets

with open(savefile, 'wb') as fout:
    pickle.dump(results, fout)

```

Train Epoch: 1 [0/864 (0%)] Loss: 1.642396
 Train Epoch: 1 [640/864 (71%)] Loss: 1.689561
 Test set: Average loss: 1.6843, Accuracy: 225/864 (26%)

Test set: Average loss: 1.5443, Accuracy: 75/216 (35%)

Train Epoch: 2 [0/864 (0%)] Loss: 1.575958
 Train Epoch: 2 [640/864 (71%)] Loss: 1.489225
 Test set: Average loss: 1.5467, Accuracy: 259/864 (30%)

Test set: Average loss: 1.4197, Accuracy: 85/216 (39%)

Train Epoch: 3 [0/864 (0%)] Loss: 1.455667
 Train Epoch: 3 [640/864 (71%)] Loss: 1.645456
 Test set: Average loss: 1.4980, Accuracy: 280/864 (32%)

Test set: Average loss: 1.4879, Accuracy: 68/216 (31%)

Train Epoch: 4 [0/864 (0%)] Loss: 1.627146
 Train Epoch: 4 [640/864 (71%)] Loss: 1.333642
 Test set: Average loss: 1.5157, Accuracy: 291/864 (34%)

Test set: Average loss: 1.4738, Accuracy: 77/216 (36%)

Train Epoch: 5 [0/864 (0%)] Loss: 1.366408
 Train Epoch: 5 [640/864 (71%)] Loss: 1.620814
 Test set: Average loss: 1.4954, Accuracy: 308/864 (36%)

Test set: Average loss: 1.5680, Accuracy: 52/216 (24%)

Train Epoch: 6 [0/864 (0%)] Loss: 1.479855
 Train Epoch: 6 [640/864 (71%)] Loss: 1.524231
 Test set: Average loss: 1.5308, Accuracy: 232/864 (27%)

Test set: Average loss: 1.5924, Accuracy: 69/216 (32%)

Train Epoch: 7 [0/864 (0%)] Loss: 1.452423

Train Epoch: 7 [640/864 (71%)] Loss: 1.479249

Test set: Average loss: 1.4871, Accuracy: 282/864 (33%)

Test set: Average loss: 1.4564, Accuracy: 68/216 (31%)

Train Epoch: 8 [0/864 (0%)] Loss: 1.562650

Train Epoch: 8 [640/864 (71%)] Loss: 1.366179

Test set: Average loss: 1.4463, Accuracy: 292/864 (34%)

Test set: Average loss: 1.5645, Accuracy: 57/216 (26%)

Train Epoch: 9 [0/864 (0%)] Loss: 1.571358

Train Epoch: 9 [640/864 (71%)] Loss: 1.383995

Test set: Average loss: 1.5485, Accuracy: 262/864 (30%)

Test set: Average loss: 1.4434, Accuracy: 83/216 (38%)

Train Epoch: 10 [0/864 (0%)] Loss: 1.512825

Train Epoch: 10 [640/864 (71%)] Loss: 1.397849

Test set: Average loss: 1.4526, Accuracy: 310/864 (36%)

Test set: Average loss: 1.4392, Accuracy: 69/216 (32%)

Train Epoch: 11 [0/864 (0%)] Loss: 1.504065

Train Epoch: 11 [640/864 (71%)] Loss: 1.454877

Test set: Average loss: 1.4286, Accuracy: 316/864 (37%)

Test set: Average loss: 1.4023, Accuracy: 84/216 (39%)

Train Epoch: 12 [0/864 (0%)] Loss: 1.484661

Train Epoch: 12 [640/864 (71%)] Loss: 1.577912

Test set: Average loss: 1.4857, Accuracy: 311/864 (36%)

Test set: Average loss: 1.5601, Accuracy: 59/216 (27%)

Train Epoch: 13 [0/864 (0%)] Loss: 1.308768

Train Epoch: 13 [640/864 (71%)] Loss: 1.430030

Test set: Average loss: 1.4538, Accuracy: 314/864 (36%)

Test set: Average loss: 1.4391, Accuracy: 68/216 (31%)

Train Epoch: 14 [0/864 (0%)] Loss: 1.500610

Train Epoch: 14 [640/864 (71%)] Loss: 1.323643

Test set: Average loss: 1.4664, Accuracy: 317/864 (37%)

Test set: Average loss: 1.4491, Accuracy: 73/216 (34%)

Train Epoch: 15 [0/864 (0%)] Loss: 1.505487

Train Epoch: 15 [640/864 (71%)] Loss: 1.498564

Test set: Average loss: 1.4050, Accuracy: 342/864 (40%)

Test set: Average loss: 1.4328, Accuracy: 82/216 (38%)

Train Epoch: 16 [0/864 (0%)] Loss: 1.382088

Train Epoch: 16 [640/864 (71%)] Loss: 1.352231

Test set: Average loss: 1.3497, Accuracy: 354/864 (41%)

Test set: Average loss: 1.3772, Accuracy: 81/216 (38%)

Train Epoch: 17 [0/864 (0%)] Loss: 1.444847

Train Epoch: 17 [640/864 (71%)] Loss: 1.479589

Test set: Average loss: 1.4247, Accuracy: 310/864 (36%)

Test set: Average loss: 1.3995, Accuracy: 72/216 (33%)

Train Epoch: 18 [0/864 (0%)] Loss: 1.405810

Train Epoch: 18 [640/864 (71%)] Loss: 1.365805

Test set: Average loss: 1.4087, Accuracy: 328/864 (38%)

Test set: Average loss: 1.4100, Accuracy: 70/216 (32%)

Train Epoch: 19 [0/864 (0%)] Loss: 1.353899

Train Epoch: 19 [640/864 (71%)] Loss: 1.427612

Test set: Average loss: 1.3424, Accuracy: 352/864 (41%)

Test set: Average loss: 1.4080, Accuracy: 73/216 (34%)

Train Epoch: 20 [0/864 (0%)] Loss: 1.324723

Train Epoch: 20 [640/864 (71%)] Loss: 1.388074

Test set: Average loss: 1.3632, Accuracy: 370/864 (43%)

Test set: Average loss: 1.4279, Accuracy: 70/216 (32%)

Train Epoch: 21 [0/864 (0%)] Loss: 1.321498

Train Epoch: 21 [640/864 (71%)] Loss: 1.404276

Test set: Average loss: 1.3650, Accuracy: 346/864 (40%)

Test set: Average loss: 1.4162, Accuracy: 73/216 (34%)

Train Epoch: 22 [0/864 (0%)] Loss: 1.354182

Train Epoch: 22 [640/864 (71%)] Loss: 1.349026

Test set: Average loss: 1.3790, Accuracy: 340/864 (39%)

Test set: Average loss: 1.4206, Accuracy: 70/216 (32%)

Train Epoch: 23 [0/864 (0%)] Loss: 1.438953

Train Epoch: 23 [640/864 (71%)] Loss: 1.404784

Test set: Average loss: 1.3888, Accuracy: 343/864 (40%)

Test set: Average loss: 1.3836, Accuracy: 76/216 (35%)

Train Epoch: 24 [0/864 (0%)] Loss: 1.476820

Train Epoch: 24 [640/864 (71%)] Loss: 1.385354

Test set: Average loss: 1.3675, Accuracy: 334/864 (39%)

Test set: Average loss: 1.3429, Accuracy: 91/216 (42%)

Train Epoch: 25 [0/864 (0%)] Loss: 1.414404

Train Epoch: 25 [640/864 (71%)] Loss: 1.376806

Test set: Average loss: 1.3799, Accuracy: 343/864 (40%)

Test set: Average loss: 1.3628, Accuracy: 85/216 (39%)

Train Epoch: 26 [0/864 (0%)] Loss: 1.311600

Train Epoch: 26 [640/864 (71%)] Loss: 1.357656

Test set: Average loss: 1.3730, Accuracy: 361/864 (42%)

Test set: Average loss: 1.3927, Accuracy: 76/216 (35%)

Train Epoch: 27 [0/864 (0%)] Loss: 1.414046

Train Epoch: 27 [640/864 (71%)] Loss: 1.416384

Test set: Average loss: 1.3482, Accuracy: 355/864 (41%)

Test set: Average loss: 1.4153, Accuracy: 80/216 (37%)

Train Epoch: 28 [0/864 (0%)] Loss: 1.341959

Train Epoch: 28 [640/864 (71%)] Loss: 1.448782

Test set: Average loss: 1.3358, Accuracy: 365/864 (42%)

Test set: Average loss: 1.3772, Accuracy: 80/216 (37%)

Train Epoch: 29 [0/864 (0%)] Loss: 1.495070

Train Epoch: 29 [640/864 (71%)] Loss: 1.516291

Test set: Average loss: 1.3247, Accuracy: 369/864 (43%)

Test set: Average loss: 1.3605, Accuracy: 82/216 (38%)

Train Epoch: 30 [0/864 (0%)] Loss: 1.385395

Train Epoch: 30 [640/864 (71%)] Loss: 1.557401

Test set: Average loss: 1.3435, Accuracy: 358/864 (41%)

Test set: Average loss: 1.3594, Accuracy: 83/216 (38%)

Train Epoch: 31 [0/864 (0%)] Loss: 1.446923

Train Epoch: 31 [640/864 (71%)] Loss: 1.305975

Test set: Average loss: 1.3675, Accuracy: 363/864 (42%)

Test set: Average loss: 1.3794, Accuracy: 81/216 (38%)

Train Epoch: 32 [0/864 (0%)] Loss: 1.323347

Train Epoch: 32 [640/864 (71%)] Loss: 1.453192

Test set: Average loss: 1.3474, Accuracy: 365/864 (42%)

Test set: Average loss: 1.3761, Accuracy: 83/216 (38%)

Train Epoch: 33 [0/864 (0%)] Loss: 1.375228

Train Epoch: 33 [640/864 (71%)] Loss: 1.427434

Test set: Average loss: 1.3464, Accuracy: 350/864 (41%)

Test set: Average loss: 1.3669, Accuracy: 85/216 (39%)

Train Epoch: 34 [0/864 (0%)] Loss: 1.444740

Train Epoch: 34 [640/864 (71%)] Loss: 1.397272

Test set: Average loss: 1.3517, Accuracy: 368/864 (43%)

Test set: Average loss: 1.3729, Accuracy: 83/216 (38%)

Train Epoch: 35 [0/864 (0%)] Loss: 1.380625

Train Epoch: 35 [640/864 (71%)] Loss: 1.420936

Test set: Average loss: 1.3412, Accuracy: 372/864 (43%)

Test set: Average loss: 1.3731, Accuracy: 82/216 (38%)

Train Epoch: 36 [0/864 (0%)] Loss: 1.448111

Train Epoch: 36 [640/864 (71%)] Loss: 1.549324

Test set: Average loss: 1.3560, Accuracy: 367/864 (42%)

Test set: Average loss: 1.3865, Accuracy: 79/216 (37%)

Train Epoch: 37 [0/864 (0%)] Loss: 1.383034

Train Epoch: 37 [640/864 (71%)] Loss: 1.386875

Test set: Average loss: 1.3549, Accuracy: 354/864 (41%)

Test set: Average loss: 1.3658, Accuracy: 82/216 (38%)

Train Epoch: 38 [0/864 (0%)] Loss: 1.380143

Train Epoch: 38 [640/864 (71%)] Loss: 1.362363

Test set: Average loss: 1.3644, Accuracy: 347/864 (40%)

Test set: Average loss: 1.3574, Accuracy: 83/216 (38%)

Train Epoch: 39 [0/864 (0%)] Loss: 1.351373

Train Epoch: 39 [640/864 (71%)] Loss: 1.429460

Test set: Average loss: 1.3293, Accuracy: 367/864 (42%)

Test set: Average loss: 1.3520, Accuracy: 87/216 (40%)

Train Epoch: 40 [0/864 (0%)] Loss: 1.384653

Train Epoch: 40 [640/864 (71%)] Loss: 1.388518

Test set: Average loss: 1.3726, Accuracy: 337/864 (39%)

Test set: Average loss: 1.3623, Accuracy: 82/216 (38%)

Train Epoch: 41 [0/864 (0%)] Loss: 1.333001

Train Epoch: 41 [640/864 (71%)] Loss: 1.432841

Test set: Average loss: 1.3746, Accuracy: 343/864 (40%)

Test set: Average loss: 1.3582, Accuracy: 83/216 (38%)

Train Epoch: 42 [0/864 (0%)] Loss: 1.336418

Train Epoch: 42 [640/864 (71%)] Loss: 1.387215

Test set: Average loss: 1.3458, Accuracy: 348/864 (40%)

Test set: Average loss: 1.3611, Accuracy: 83/216 (38%)

Train Epoch: 43 [0/864 (0%)] Loss: 1.293901

Train Epoch: 43 [640/864 (71%)] Loss: 1.370599

Test set: Average loss: 1.3525, Accuracy: 373/864 (43%)

Test set: Average loss: 1.3693, Accuracy: 82/216 (38%)

Train Epoch: 44 [0/864 (0%)] Loss: 1.466611

Train Epoch: 44 [640/864 (71%)] Loss: 1.460432

Test set: Average loss: 1.3296, Accuracy: 351/864 (41%)

Test set: Average loss: 1.3510, Accuracy: 88/216 (41%)

Train Epoch: 45 [0/864 (0%)] Loss: 1.367427

Train Epoch: 45 [640/864 (71%)] Loss: 1.259825

Test set: Average loss: 1.3891, Accuracy: 340/864 (39%)

Test set: Average loss: 1.3638, Accuracy: 84/216 (39%)

Train Epoch: 46 [0/864 (0%)] Loss: 1.466855

Train Epoch: 46 [640/864 (71%)] Loss: 1.321473

Test set: Average loss: 1.3629, Accuracy: 347/864 (40%)

```

Test set: Average loss: 1.3608, Accuracy: 84/216 (39%)

Train Epoch: 47 [0/864 (0%)]    Loss: 1.357319
Train Epoch: 47 [640/864 (71%)] Loss: 1.382298
Test set: Average loss: 1.3783, Accuracy: 346/864 (40%)

Test set: Average loss: 1.3587, Accuracy: 83/216 (38%)

Train Epoch: 48 [0/864 (0%)]    Loss: 1.392995
Train Epoch: 48 [640/864 (71%)] Loss: 1.438382
Test set: Average loss: 1.3089, Accuracy: 385/864 (45%)

Test set: Average loss: 1.3561, Accuracy: 85/216 (39%)

Train Epoch: 49 [0/864 (0%)]    Loss: 1.334245
Train Epoch: 49 [640/864 (71%)] Loss: 1.312263
Test set: Average loss: 1.3657, Accuracy: 333/864 (39%)

Test set: Average loss: 1.3622, Accuracy: 84/216 (39%)

```

```
[28]: import seaborn as sns
from sklearn.metrics import confusion_matrix, recall_score

fig, (ax1, ax2, ax3) = plt.subplots(1, 3, figsize=(18, 5))

# Liste pour stocker les résultats du dernier modèle chargé
last_results = None
last_label = None

for filename in os.listdir(savedir):
    if filename.endswith('.pkl'):
        with open(os.path.join(savedir, filename), 'rb') as fin:
            results = pickle.load(fin)
            label = filename[:-4] # nom sans .pkl
            last_results = results
            last_label = label

            # --- Courbes de LOSS (ax1) ---
            ax1.plot(results['train_loss'], '--', label=f'{label} train')
            ax1.plot(results['val_loss'], '-', label=f'{label} val')
            ax1.set_ylabel('Loss (MSE)') # Mise à jour du label Y pour MSE
            ax1.set_xlabel('epochs')
            ax1.set_title('Train vs Validation Loss')

            # --- Courbes d'ACCURACY (ax2) ---
            ax2.plot(results['train_accuracy'], '--', label=f'{label} train')
            ax2.plot(results['val_accuracy'], '-', label=f'{label} val')
            ax2.set_ylabel('Accuracy')
            ax2.set_xlabel('epochs')
            ax2.set_title('Train vs Validation Accuracy')

            # --- Courbes de RECALL (ax3) ---
            ax3.plot(results['train_recall'], '--', label=f'{label} train')
            ax3.plot(results['val_recall'], '-', label=f'{label} val')
            ax3.set_ylabel('Recall')
            ax3.set_xlabel('epochs')
            ax3.set_title('Train vs Validation Recall')

plt.tight_layout()
plt.show()
```

```

        ax2.plot(results['train_acc'], '--', label=f'{label} train')
        ax2.plot(results['val_acc'], '-', label=f'{label} val')

    # AJOUT DU MACRO RECALL SI DÉFINI
    if 'val_macro_recall' in results:
        ax2.plot(results['val_macro_recall'], '-.', label=f'{label} Macro Recall')

    ax2.set_ylabel('Accuracy / Macro Recall')
    ax2.set_xlabel('epochs')
    ax2.set_title('Train vs Validation Accuracy')

# --- MATRICE DE CONFUSION (ax3) ---
# Nécessite les prédictions et targets finales (non incluses dans le .pkl standard)
# SI vous avez sauvégarde all_val_preds et all_val_targets dans results:

print("Last results keys:", last_results.keys() if last_results else "No results loaded")

if last_results and 'final_val_preds' in last_results:
    y_true = np.array(last_results['final_val_targets'])
    y_pred = np.array(last_results['final_val_preds'])

    cm = confusion_matrix(y_true, y_pred)

    # Affichage de la Heatmap
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', ax=ax3, cbar=False)
    ax3.set_title(f'Matrice de Confusion ({last_label})')
    ax3.set_xlabel('Prédiction')
    ax3.set_ylabel('Vraie Classe')

# Légende pour ax1 et ax2
ax1.legend()
ax2.legend()

plt.tight_layout()
plt.show()

```

Last results keys: dict_keys(['name', 'lr', 'train_loss', 'train_acc', 'val_loss', 'val_acc'])

